

# Week 00

## Building Production Ready ML Systems

A comprehensive 10-week journey to set the ground for transformation of you from a model builder to a E2E ML engineer



# About Your Instructors

## **Isuru Alagiyawanna**

Lead Instructor

7+ years experience in AI research and Academics

4+ years experience in industry applications



# Welcome to Day 01: Course Introduction

This course bridges the gap between academic ML knowledge and industry-ready ML engineering skills

Today you'll learn:

- The structure of our 10-week journey
- Required tools, technologies, and resources
- Project expectations and support channels



# Today's Agenda

## Foundation (30 min)

- Course introduction
- Module-by-module summary
- Required software

## Implementation (20 min)

- Tools & libraries
- Mini projects overview
- Capstone project intro

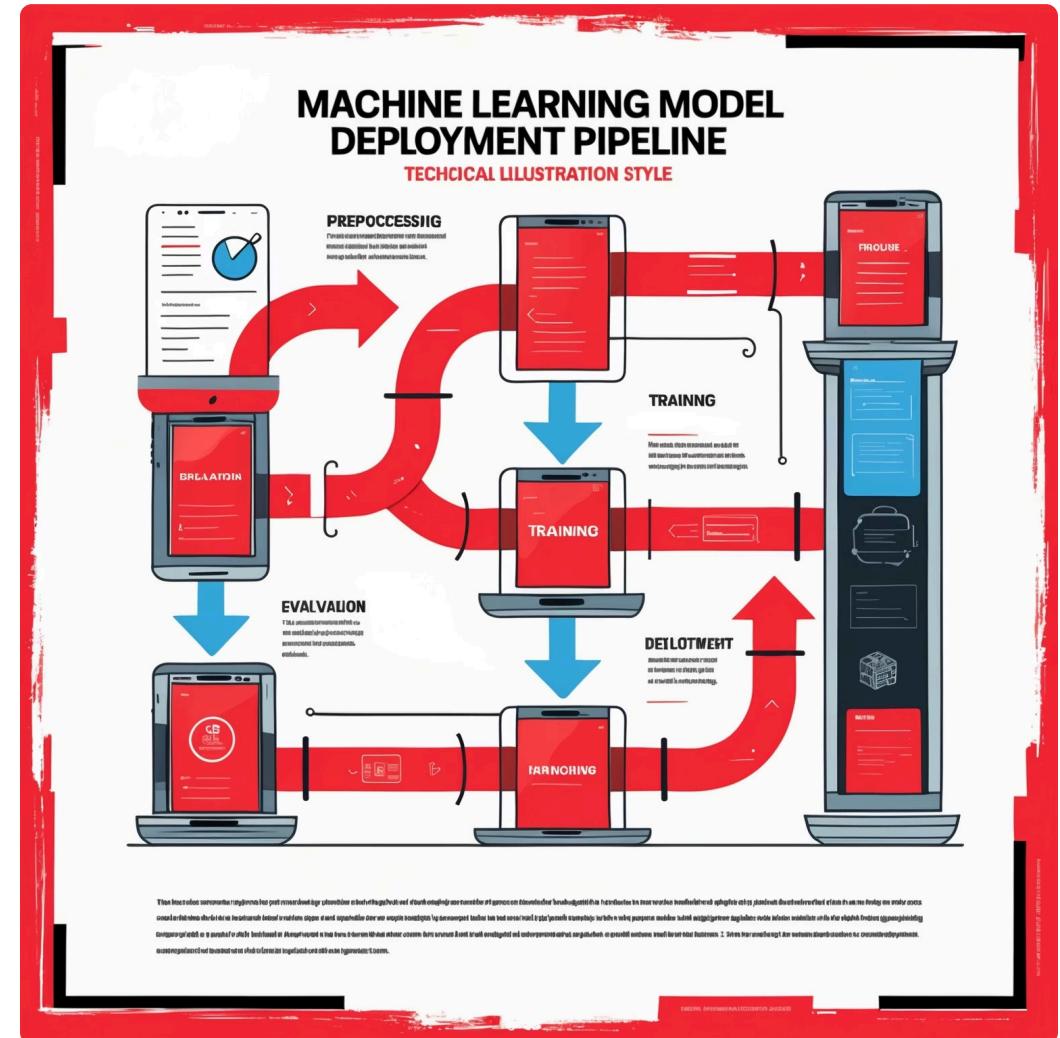
## Operations (10 min)

- Hackathon details
- Assessment criteria
- Communication channels

# Why "Production-Ready" Matters

In industry, ML excellence requires more than accurate models:

- Deployable in production environments
- Scalable to handle real-world data volumes
- Maintainable through infrastructure changes
- Monitorable for performance degradation
- Reproducible for debugging and auditing



# The ML Production Gap

## Academic ML

Focus on algorithm understanding  
and model accuracy on static  
datasets

## Our Approach

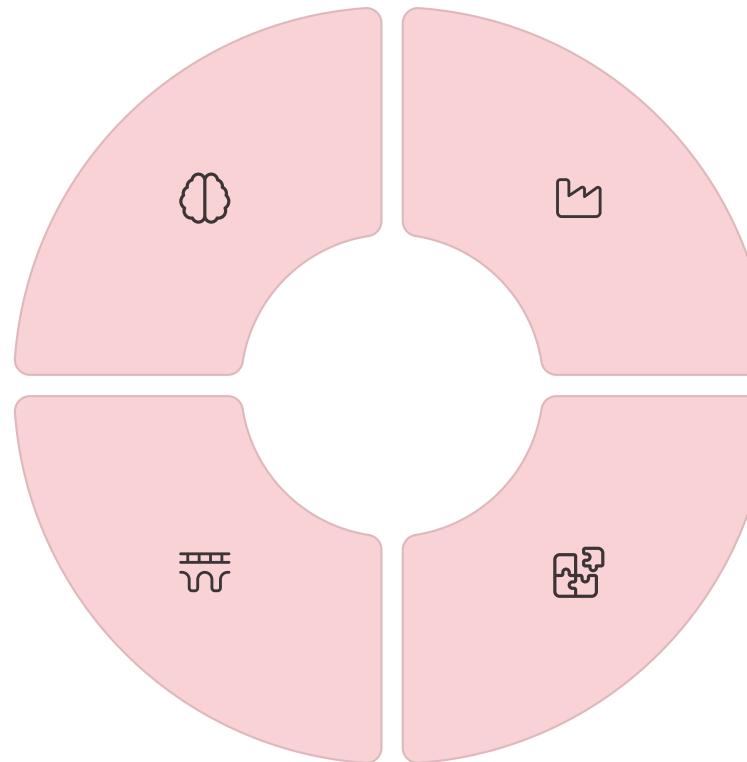
Practical projects using industry tools  
to build and deploy end-to-end ML  
systems

## Industry ML

Emphasis on reliable systems that  
handle changing data and scale with  
demand

## The Gap

Skills in containerization,  
orchestration, monitoring, and  
scalable processing



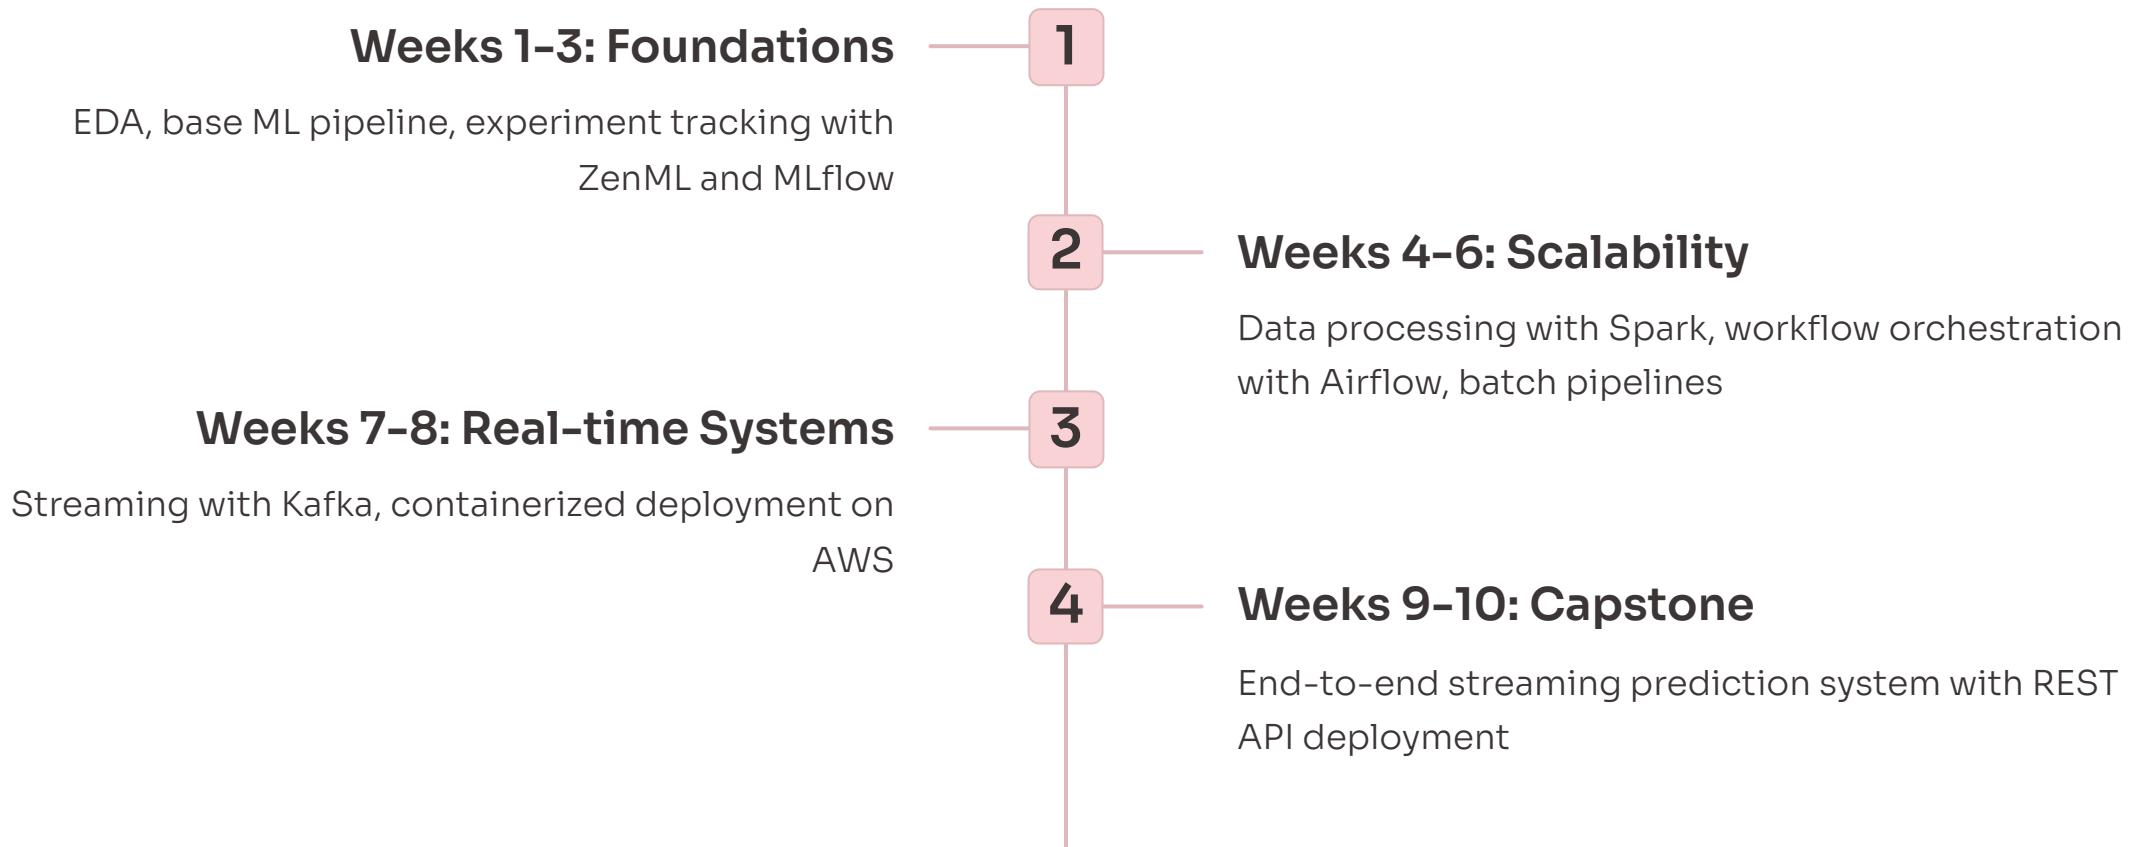
# Zuu Crew's Teaching Philosophy

## Industry-Aligned Learning

- Learn by building: hands-on mini-projects
- Real-world tools: Kafka, Spark, Airflow, MLflow
- Practical workflows: CI/CD, monitoring, testing
- Failure resilience: error handling, fallbacks



# Your Learning Journey: 10-Week Roadmap

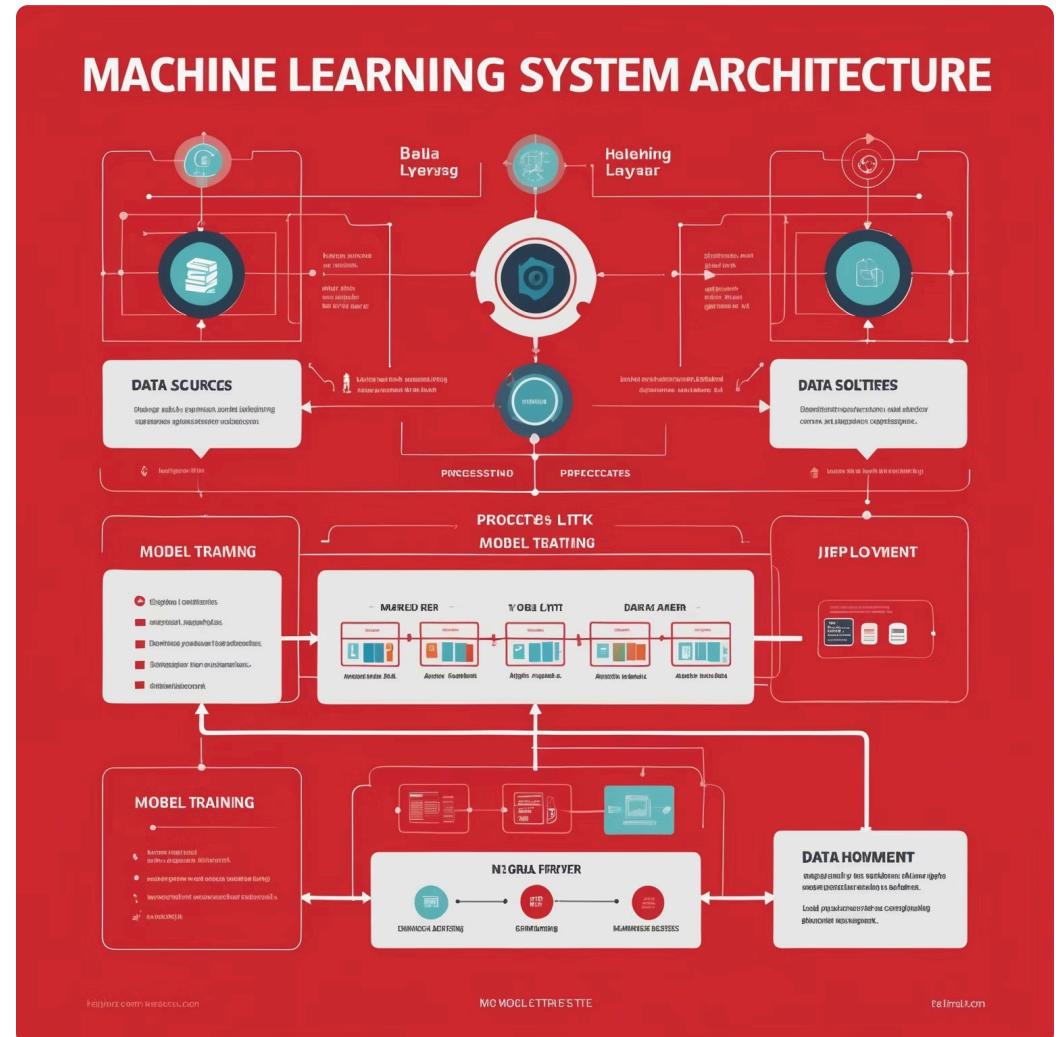


# Week 1: Capstone Architecture + EDA

## Learning Objectives:

- Understand end-to-end ML system architecture
- Plan component interactions and data flow
- Perform exploratory data analysis on the capstone dataset
- Define feature engineering strategies

Deliverables: EDA notebook, architecture diagram, feature engineering plan



# Week 2: Base ML Pipeline

## Learning Objectives:

- Build reusable scikit-learn pipelines
- Implement feature engineering transformers
- Train and optimize XGBoost models
- Create modular, testable ML code

Deliverables: Working ML pipeline with metrics reporting

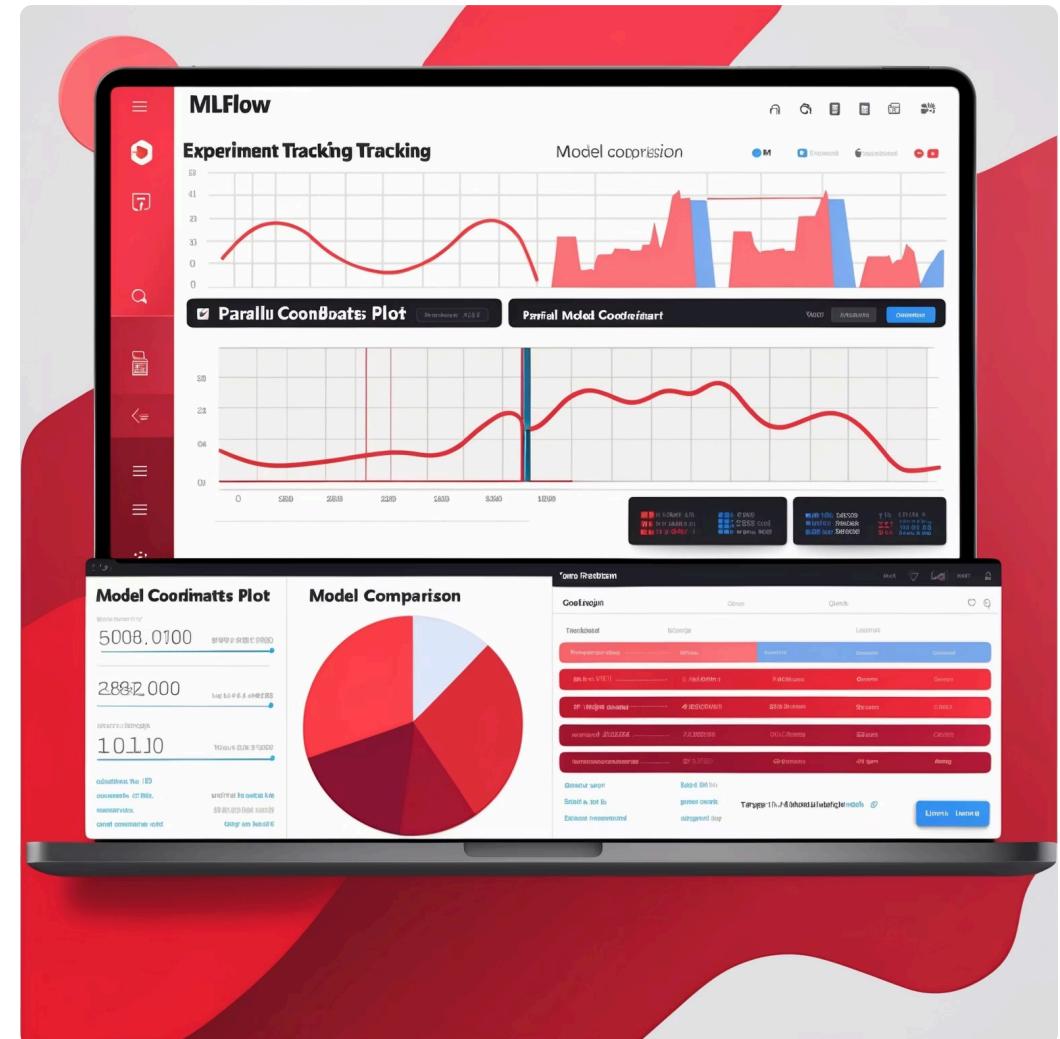
```
defines:  
    transformer = []  
    # Boost model with meorme=caumattangage en transmission state capme  
    # XGBoost model  
    transformer = [ ]  
    # remBoscLoy treem effimic Regressioh building  
    #  
    transformer = []  
    #ite otaghet Loarl gretit phoneL yohne chand mew impoing  
    # XGBoost = get[]  
    #  
    transformer model fit and entebba()  
    # da boost = (CUbcanoGd-Udection(1, n_estimators)  
    # estera() iCBoost one tchis tandomatic tewmpe  
    #  
    transformation = transform  
    # Boost model | e(yeet!  
    # XGBoost model) -> lew-inivation()  
    # XGBoost premodul emmhsutA)  
    #  
    # aogood  
    # retoast model)  
    # schueest PegoLead(Lat topime))  
    # -schueest PegoLead(Lat topime))  
    # -schueest PegoLead(Lat topime))  
    # koodgrtt  
    # geoost.  
    # Re: JX XGBoost forl (Tunavista hie ppehant)  
    # This equan ovisek etabon repesent an elating amvalant elatende ame
```

# Week 3-4: ZenML + MLflow Integration

## Learning Objectives:

- Track experiments systematically
- Version models and datasets
- Compare model performance
- Create reproducible ML pipelines

Deliverables: Mini-project 1 - Tracked experiments and registered models



# Week 5: Scalable Processing with Apache Spark

## Learning Objectives:

- Process large datasets with PySpark
- Implement distributed ML algorithms
- Optimize Spark jobs for performance
- Monitor resource utilization

Deliverables: PySpark data processing pipeline



# Week 6: Orchestration using Apache Airflow

## Learning Objectives:

- Design and implement DAGs
- Schedule recurring ML workflows
- Handle dependencies between tasks
- Monitor pipeline execution

Deliverables: Airflow DAGs for ML pipeline orchestration

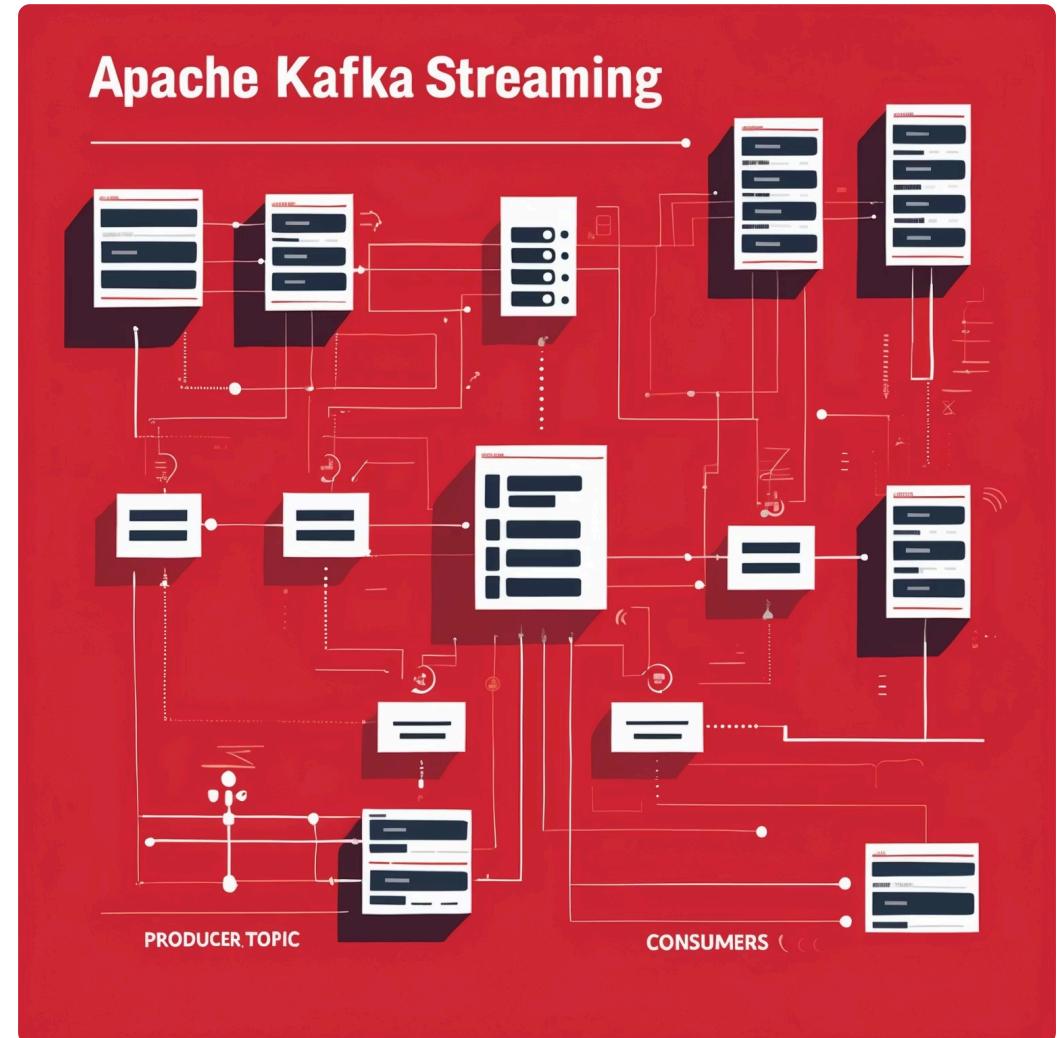


# Week 7: Real-time Pipelines with Apache Kafka

## Learning Objectives:

- Implement Kafka producers and consumers
- Process streaming data
- Design fault-tolerant pipelines
- Handle real-time prediction requests

Deliverables: Kafka streaming pipeline with real-time predictions



# Week 8: Containerized Deployment on AWS

## Learning Objectives:

- Containerize ML applications with Docker
- Configure Docker Compose for multi-service deployment
- Deploy to AWS infrastructure
- Implement monitoring and logging

Deliverables: Containerized ML system deployed on AWS



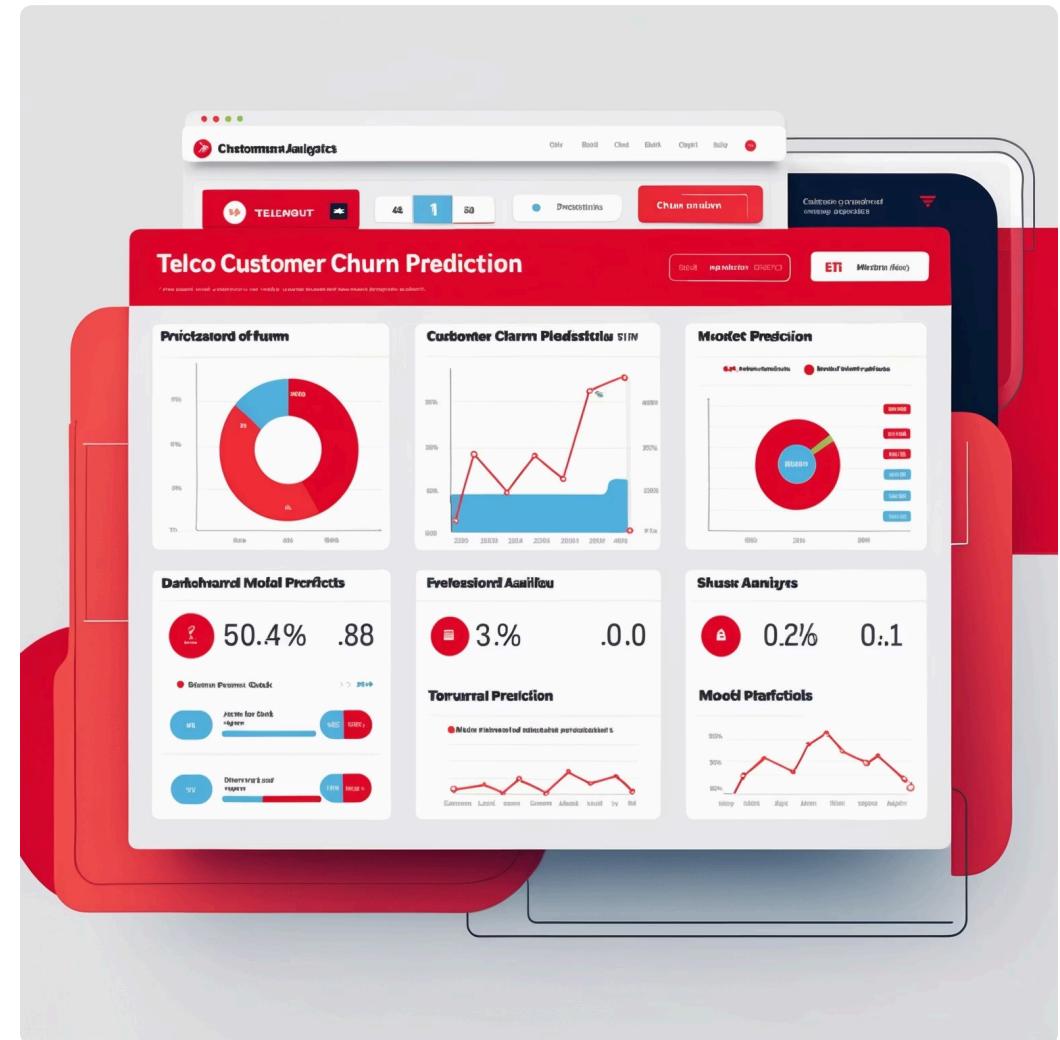
# Week 9: Mini-Project 2 - Batch Pipeline

## Project Scope:

Build a complete batch ML pipeline for Telco Customer Churn prediction

- Process data with Spark
- Orchestrate with Airflow
- Track with MLflow
- Automate retraining

Deliverables: Working end-to-end batch pipeline with documentation



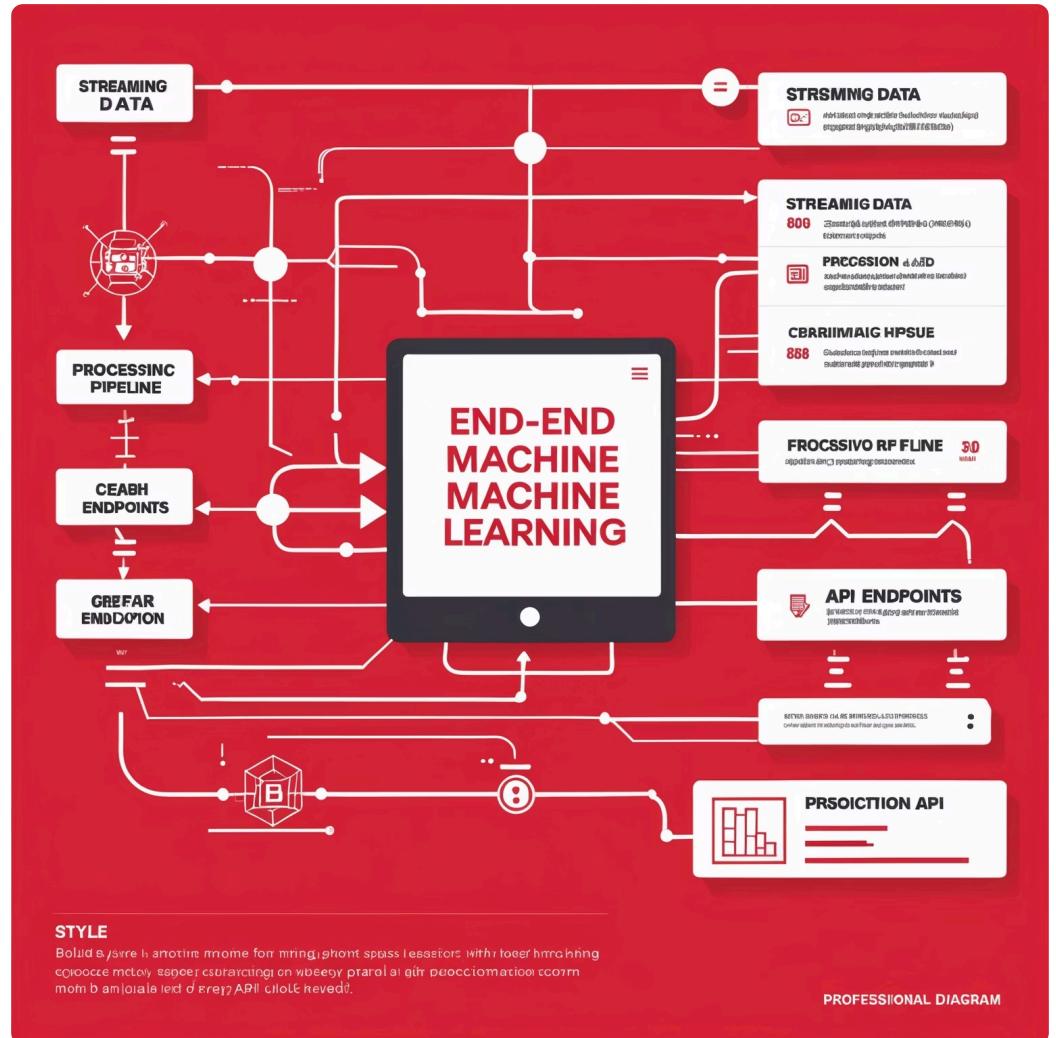
# **Weeks 10–11: Capstone Project**

## **Project Challenge:**

## Build a complete real-time streaming prediction pipeline

- Part 1: Streaming ingestion and processing
  - Part 2: Deployment and monitoring

**Deliverables:** Fully deployed REST API serving predictions with documentation and testing



# Required Software Installation



## Python 3.9+

Core programming language for all our ML implementations. We recommend Python 3.9 for optimal compatibility with all libraries.



## Docker & Docker Compose

Containerization platform for creating reproducible environments and deployable applications.

**Homework:** Complete all installations before Week 2 begins!



## Git

Version control system for tracking code changes and collaborating on projects. Required for accessing course materials.



## Java 8+

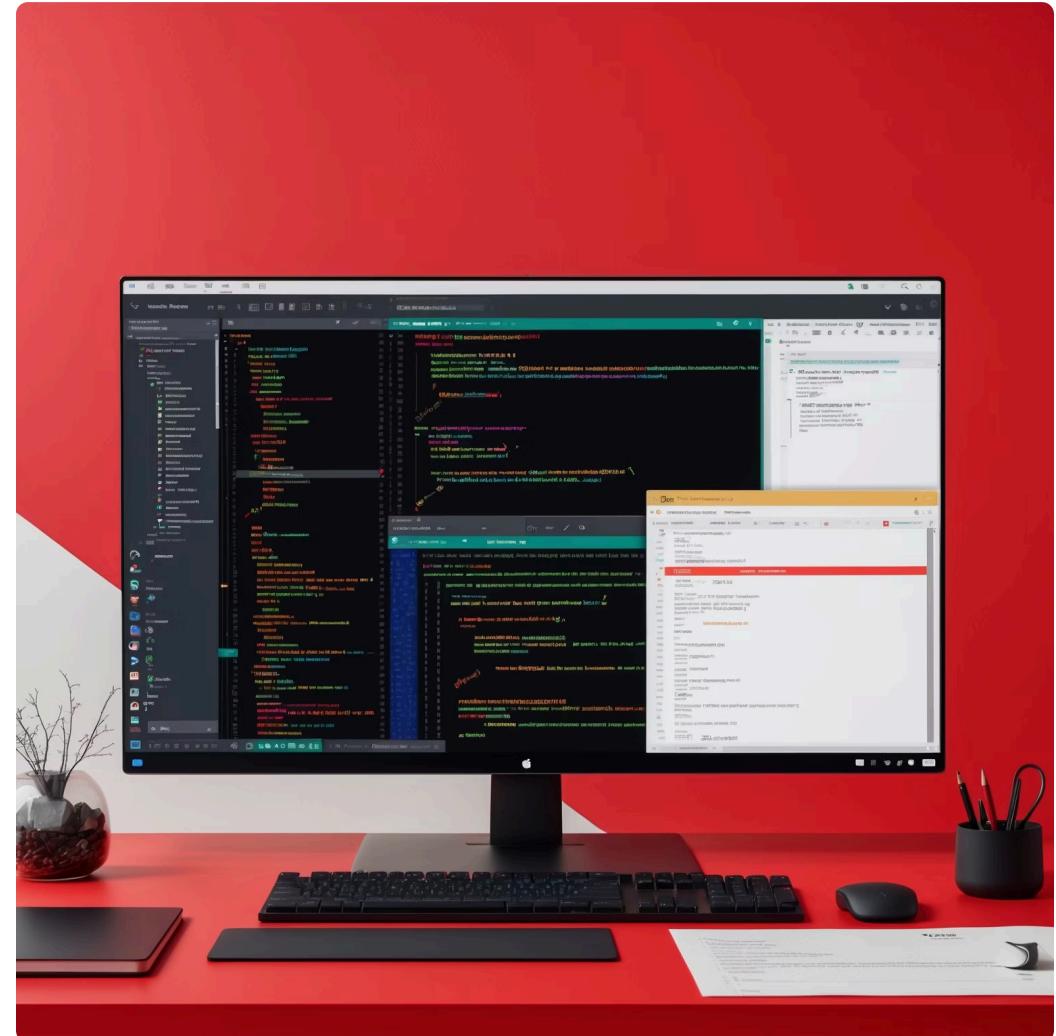
Required runtime for Apache Spark. OpenJDK 11 is recommended for best compatibility and performance.

# Development Environment Setup

## Recommended Setup:

- VSCode with Python and Docker extensions
- Jupyter Lab for interactive analysis
- Terminal access for command-line tools
- GitHub account for assignments

All course examples will assume this environment configuration for consistency.



# Environment Management

## Conda

Cross-platform package and environment manager that handles both pip and conda packages

- Creates isolated environments with specific Python versions
- Installs binary packages without compilation
- Simplifies dependency management

## Virtualenv/venv

Lightweight Python-specific virtual environment tool

- Standard in Python's standard library (venv)
- Faster creation than Conda
- Uses pip for package management

The course will use Conda for consistency, but either option is acceptable for your work.

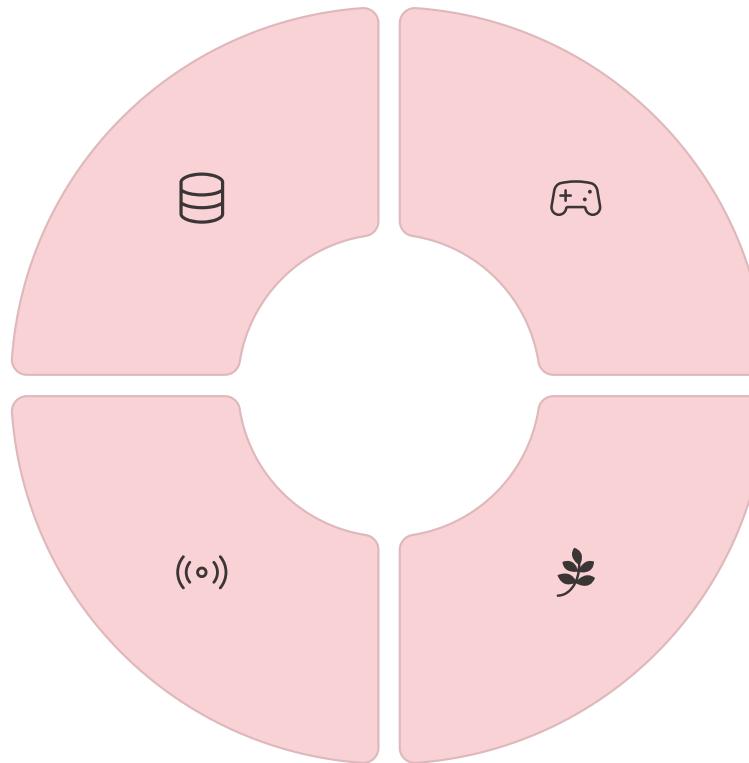
# Key Python Libraries

## Data Processing

- Pandas: tabular data manipulation
- NumPy: numerical computing
- PySpark: distributed data processing

## Infrastructure

- apache-airflow: orchestration
- kafka-python: streaming
- boto3: AWS integration



## ML Modeling

- scikit-learn: traditional ML algorithms
- XGBoost: gradient boosting
- Hyperopt: hyperparameter optimization

## ML Ops

- ZenML: ML pipeline management
- MLflow: experiment tracking
- FastAPI: model serving

# Data Processing Libraries Deep Dive

## Pandas

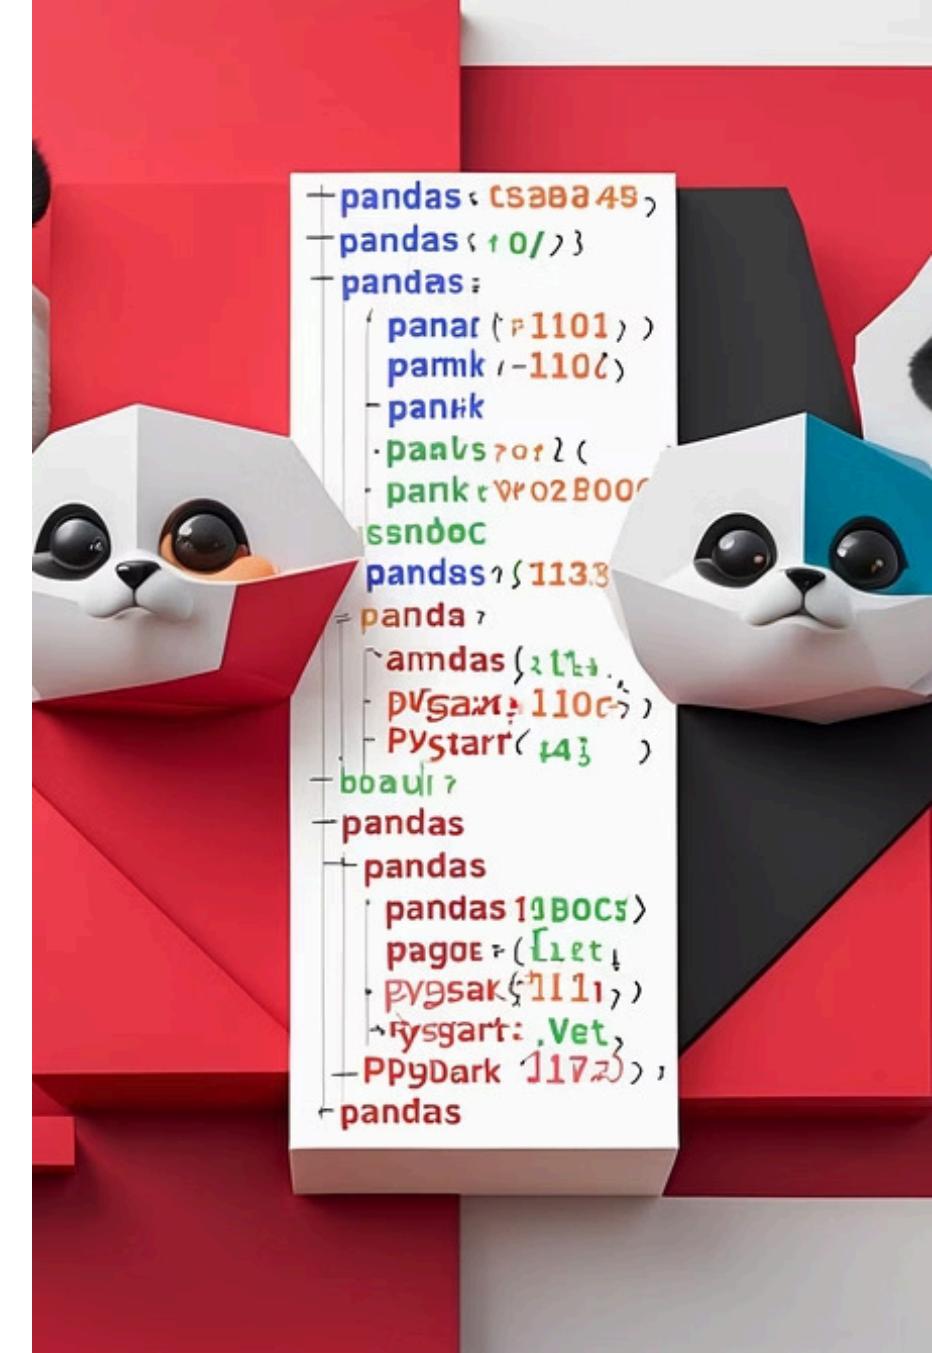
## In-memory tabular data processing:

- DataFrame and Series data structures
  - Data cleaning, transformation, and aggregation
  - Memory-optimized for smaller datasets

## PySpark

## Distributed data processing for large datasets:

- DataFrame API similar to Pandas
  - SQL query support
  - Built-in ML algorithms (MLlib)
  - Scales across clusters



# ML Tracking and Pipeline Tools



## MLflow

Open-source platform for managing the ML lifecycle:

- Tracking: log parameters, metrics, and artifacts
- Projects: package code for reproducibility
- Models: package models for deployment

## ZenML

Framework for creating reproducible ML pipelines:

- Pipeline versioning
- Artifact tracking
- Integration with MLflow
- Multi-stage workflows



# Container Tools Overview

## Docker

Container platform for packaging applications:

- Dockerfile: defines environment
- Images: immutable snapshots
- Containers: running instances
- Registry: for sharing images

## Docker Compose

Multi-container orchestration:

- YAML configuration
- Service definitions
- Networking between containers
- Volume management

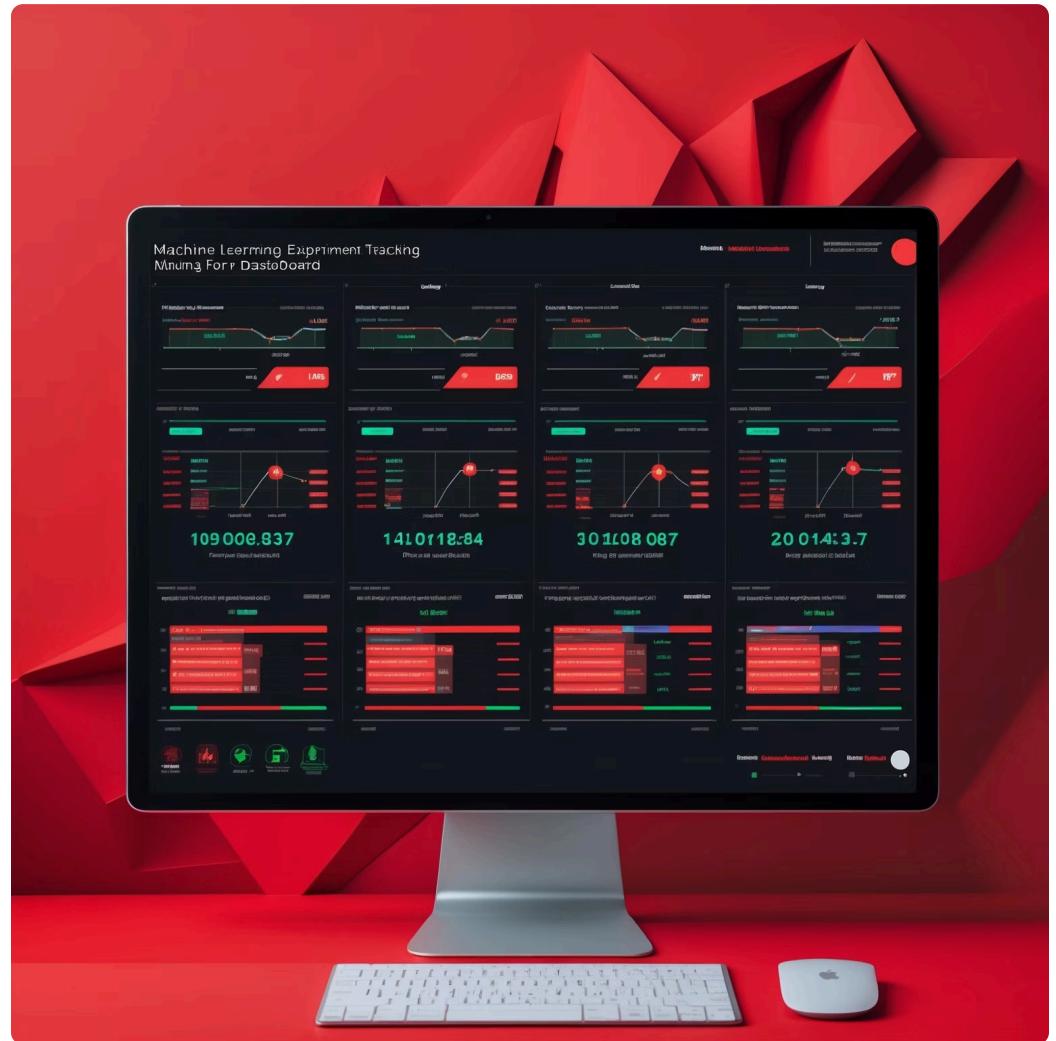
# Mini Project 1: Experiment Tracking

## Project Objective:

Implement a reproducible ML pipeline with experiment tracking

## Components:

- ZenML pipeline definition
- MLflow experiment tracking
- Model versioning and registration
- Performance comparison dashboard



# Mini Project 1: Technical Requirements

## Pipeline Structure

- Data ingestion step
- Data validation step
- Feature engineering step
- Model training step
- Evaluation step

## MLflow Integration

- Parameter logging
- Metric tracking
- Artifact storage
- Run comparison

## Deliverables

- Working ZenML pipeline
- MLflow experiment UI
- Model registry
- Documentation

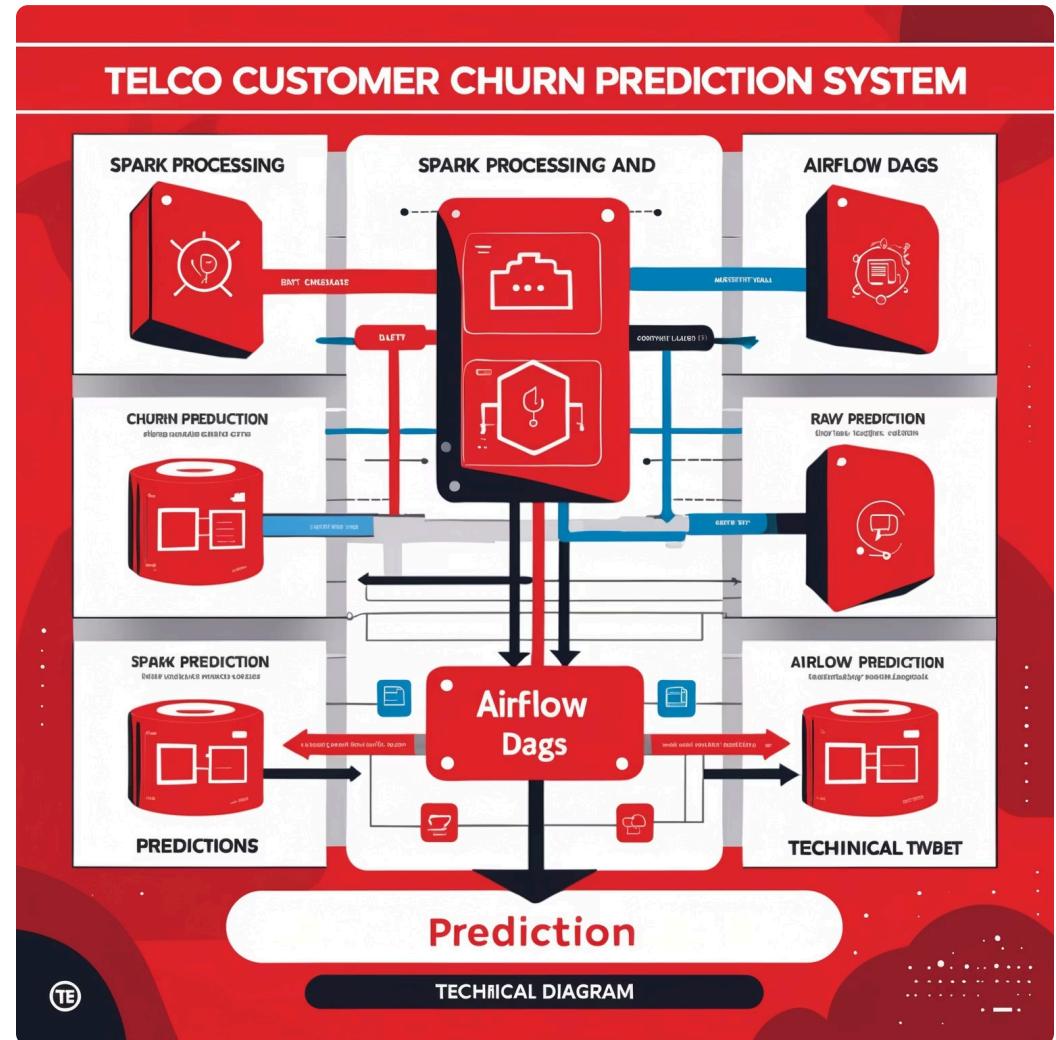
# Mini Project 2: Batch Pipeline

## **Project Objective:**

# Build an end-to-end batch prediction pipeline for Telco Customer Churn

## Components:

- Spark-based data processing
  - Airflow orchestration DAGs
  - Automated model retraining
  - Batch prediction generation



# Mini Project 2: Technical Requirements



## Data Processing

PySpark-based ETL pipeline:

- Feature engineering
- Data cleaning
- Train/test splitting

## Orchestration

Airflow DAGs for:

- Daily data ingestion
- Weekly model retraining
- Daily prediction generation

## Monitoring

Performance tracking:

- Data drift detection
- Model metrics logging
- Pipeline execution stats

# Capstone Project: Technical Requirements

## Functional Requirements:

- Ingest streaming data from Kafka
- Process data in real-time with Spark Streaming
- Make predictions using deployed ML model
- Serve predictions via REST API
- Log all operations for auditability

## Non-Functional Requirements:

- Latency < 500ms for prediction requests
- Throughput of 100+ requests/second
- 99.9% system availability
- Comprehensive error handling
- Monitoring dashboards

# Hackathon Challenge

**Test your skills in a 120-hour challenge**

- **Format:** Solo competition
- **Theme:** Solve a provided E2E ML use case
- **Challenge:** Build a E2E ML pipeline
- **Prizes:** Course credit, Zuu Crew Merch and exclusive mentorship sessions

Top 10 Candidates will present their solutions live to the class.



# Hackathon: Judging Criteria



## Technical Implementation (40%)

- Correct use of ZenML and MLflow
- Pipeline structure and modularity
- Code quality and organization



## ML Engineering Best Practices (30%)

- Data validation
- Model versioning
- Error handling
- Documentation



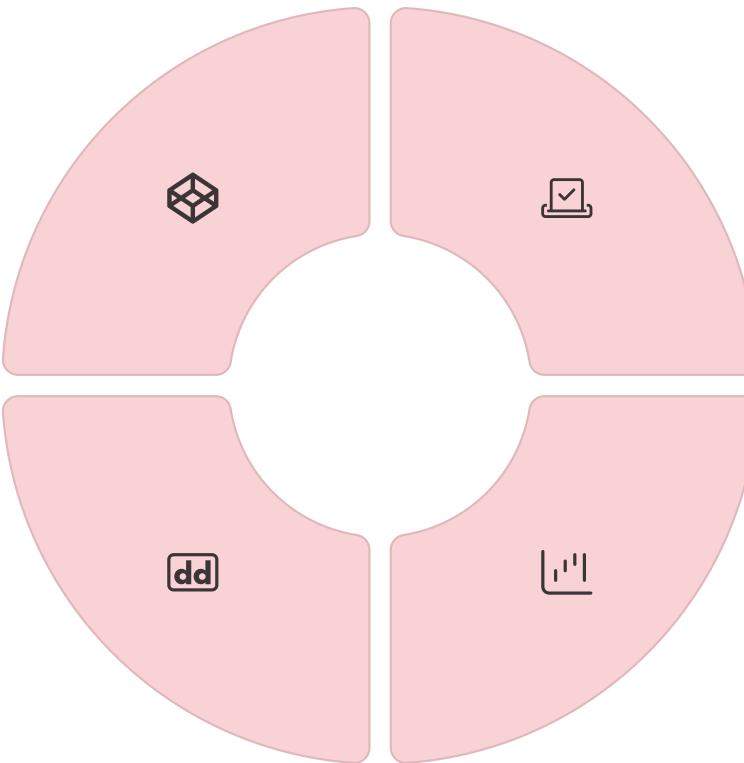
## Innovation & Presentation (30%)

- Creative approaches
- Clear explanation of design choices
- Demonstration of working pipeline

# Assessment & Grading Criteria

## Code Quality (30%)

- Readability and PEP8 compliance
- Modularity and reusability
- Comprehensive docstrings
- Thoughtful variable naming



## System Architecture (20%)

- Component integration
- Scalability considerations
- Production readiness
- Documentation completeness

## Best Practices (30%)

- Unit and integration testing
- CI-style validation checks
- Robust error handling
- Logging and monitoring

## Model Performance (20%)

- Appropriate metrics
- Validation methodology
- Baseline comparisons
- Feature importance analysis



# Communication & Support Channels

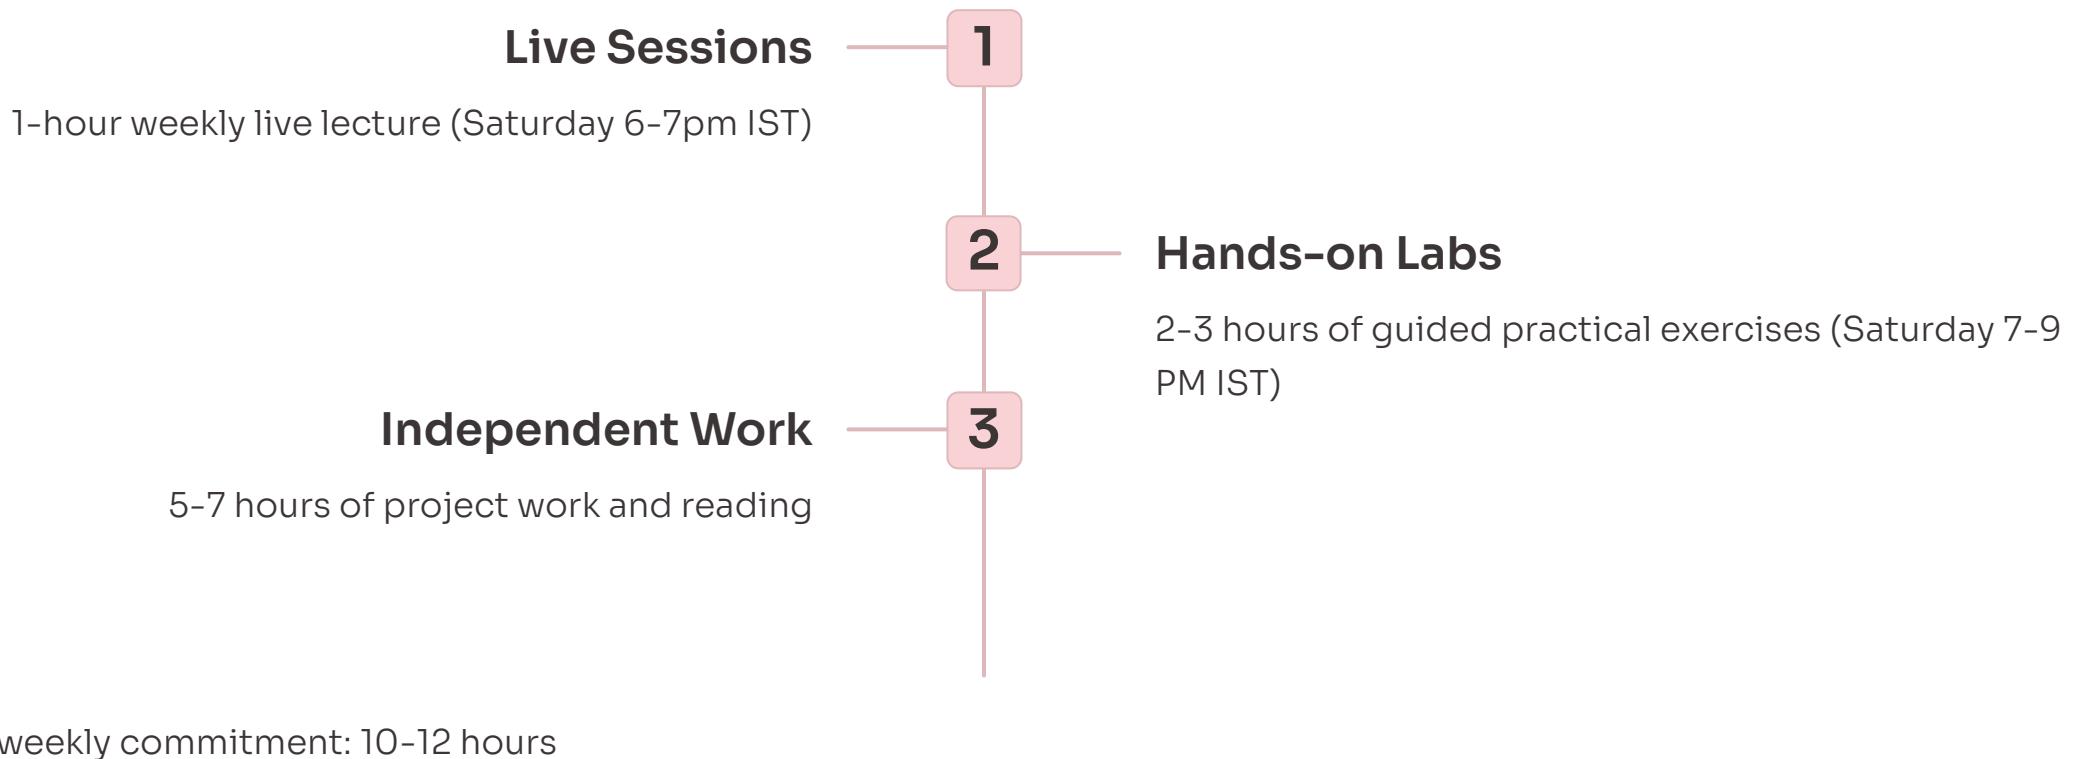
## Discord Community

- **Link:**  
<https://discord.gg/5FdQ5x5h>
- **Channels:**
  - #announcements
  - #week1, #week2, etc.
  - #help-technical
  - #general-discussion

## Office Hours & Resources

- **Office Hours:** Weekends (Sat/Sun)
- **AMA Threads:** Weekly Q&A
- **GitHub Repo:** Shared code examples
- **Response Time:** Within 24 hours

# Weekly Schedule & Commitment



# Pre-Course Preparation

## Required Tutorial Videos:

1. [Python OOP](#)
2. [FastAPI](#)
3. [Docker](#)
4. [Docker Compose](#)



# Industry Context: Why These Tools Matter



## Tool Adoption in Industry

According to the 2023 ML Tools Survey:

- 76% of companies use Docker for ML deployment
- 64% use Airflow for orchestration
- 52% use Spark for large-scale data processing
- 48% use Kafka for real-time data pipelines



## Skills Gap in the Market

LinkedIn's 2023 Emerging Jobs Report shows:

- ML Engineers command 25% higher salaries than Data Scientists
- Production ML skills are mentioned in 82% of senior ML job postings
- MLOps expertise is the fastest-growing skill requirement

# Career Opportunities After This Course

Graduates of this course have pursued various career paths in the machine learning and data engineering space:

## ML Engineer

Build and deploy machine learning models at scale in production environments

## MLOps Specialist

Focus on the operational aspects of machine learning systems including monitoring, governance, and CI/CD pipelines

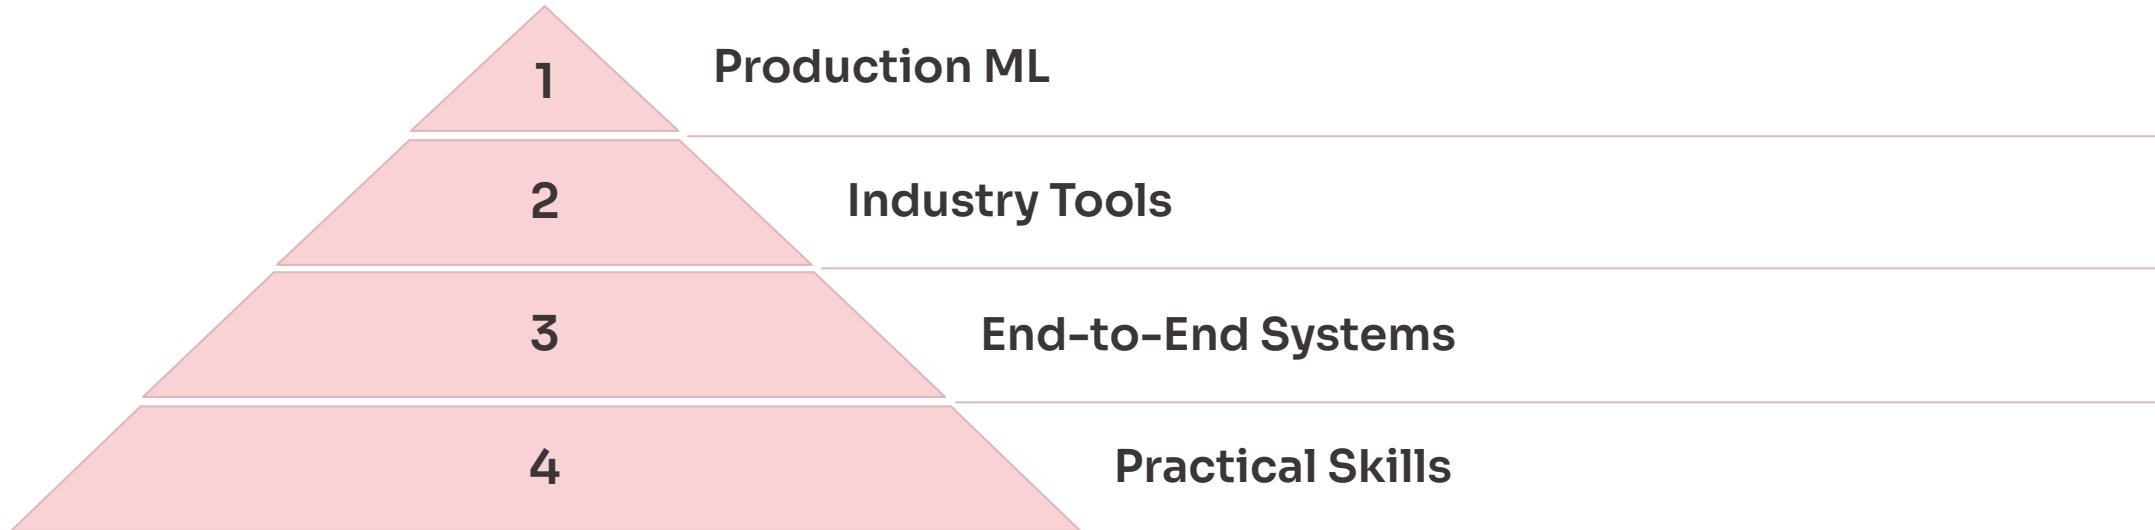
## ML Skill Specialist

Develop specialized expertise in specific ML domains like NLP, computer vision, or reinforcement learning

## ML Platform Engineer

Build internal ML platforms that enable data scientists to productionize their models

# Key Takeaways



This course bridges the gap between academic machine learning and real-world ML engineering. You'll learn to build complete, scalable, and maintainable machine learning systems using industry-standard tools and best practices.

By the end, you'll have practical experience with every component of the ML lifecycle, from data ingestion to model deployment and monitoring.

# Q & A

Any questions about the course structure or requirements?

Remember to check the pre-work instructions in the #week1 channel on Discord!

[Join Discord](#)

[Email Instructors](#)

